

neural networks in ocaml

colin shaw

07.14.2016

1. introduction

- a. welcome to computer science club
- b. thank you to revunit for sponsoring
- c. what are we talking about
 - i. neural networks
 - 1. motivation
 - 2. history
 - 3. types
 - 4. theory
 - 5. implementation
 - ii. blas / lapack
 - iii. ocaml
 - 1. applicative / functional programming
 - 2. imperative features
 - 3. interface with other languages
 - 4. modular design

2. motivations

- a. problem solving
 - i. problems we can solve explicitly
 - 1. closed form solutions
 - ii. problems we cannot solve explicitly
 - 1. do not understand the problem
 - 2. do not understand how to optimize specific solution
 - 3. desire more organic solution
 - iii. learning algorithms solve problems without us knowing how
- b. biological imperative
 - i. solving problems without understanding explicit solutions
 - ii. modeling neurons since we know they work

3. alternatives


- a. depends on type of problem
- b. some examples
 - i. regression
 - ii. analysis of variance
 - iii. k-means and similar clustering
 - iv. principal component analysis
 - v. support vector machine
 - vi. self organizing maps

4. brief history

- a. 1800s
 - i. linear regression (legendre, gauss)
- b. 1900s
 - i. gradient descent (hadamard)
- c. 1940s
 - i. early architecture (mcculloch, pitts)
- d. 1950s
 - i. simple supervised learning (perceptron; rosenblatt)
 - ii. visual cortex experiments (hubel, wiesel)
- e. 1960s
 - i. unsupervised learning
 - ii. gradient descent backpropagation
 - iii. multiple layer perceptron (deep learning)
 - iv. grossberg, ivakhnenko
- f. 1970s
 - i. convolution, subsampling (fukushima)
- g. 1980s
 - i. convolutional backpropagation (rumelhart)
 - ii. hopfield network, boltzmann machines
- h. 1990s
 - i. convergence optimization
 - ii. neural networks start winning pattern recognition contests
- i. 2000s
 - i. gpu become prevalent in non-graphical computation
- j. 2010s
 - i. mnist record broken by neural network
 - ii. gpu-based neural network surpasses human vision recognition
 - iii. neural network beats world class go player
 - iv. neural network drives a car

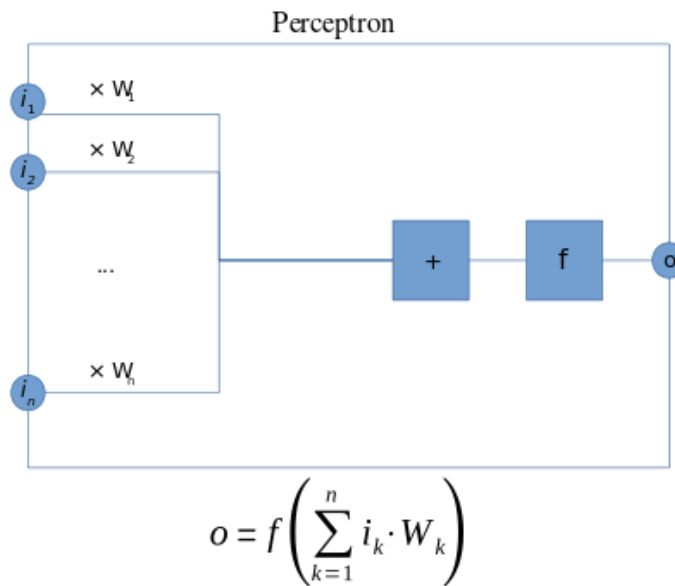
5. types of neural network

a. the inspiration




b. perceptron


- 1957 by rosenblatt as mark i perceptron (physical machine)
- sparked hype about artificial intelligence
- can only solve linearly separable problems
- cannot represent xor




- c. multiple layer perceptron
 - i. one or more hidden layers
 - ii. can represent xor



- d. convolutional
 - i. term derived generally from convolution operations on images
 - ii. preliminary feature detectors
 - iii. typically followed by mlp hidden layers



- e. recurrent
 - i. can learn based on training order
 - ii. must abide shannon sampling theorem
 - iii. scaling can be difficult
 - iv. possible implementation
 - 1. inter-layer recurrence
 - 2. delay line recurrence



6. how the multiple layer perceptron works

- a. simple regression example
 - i. $y = m \cdot x + b$
 - ii. compute output error for given input
 - 1. $\Delta y = y_{\text{new}} - y_{\text{old}}$
 - iii. compute derivatives
 - 1. $\frac{dy}{dm} = x$
 - 2. $\frac{dy}{db} = 1$
 - iv. update a and b
 - 1. let η be a learning constant
 - 2. $\Delta m = \eta \cdot \frac{dy}{dm} \cdot \Delta y = \eta \cdot x \cdot \Delta y$
 - 3. $\Delta b = \eta \cdot \frac{dy}{db} \cdot \Delta y = \eta \cdot \Delta y$

b. cost function


- i. a measure of output error
- ii. in our regression example: $\Delta y = y_{\text{new}} - y_{\text{old}}$
 1. the gradient of a cost function (∇C)
 2. cost function itself is $\frac{1}{2} (y_{\text{new}} - y_{\text{old}})^2$
- iii. in general can be any analytic function
- iv. to be able to perform mini-batch optimizations
 1. must satisfy $C = \frac{1}{n} \sum_{x=1}^n C_x$

c. activation function

- i. purpose
 1. facilitating better convergence and problem conditioning
 2. avoid stalling for zero-valued weights
 3. avoid excessive change for large weights and errors
- ii. common examples
 1. hidden layer (zero-centered, bounded)
 - a. tanh
 - b. softsign
 2. output layer (non-negative, bounded)
 - a. rectified linear (relu)
 - b. softplus

c. putting it all together

i. Network diagram



ii. compute forward pass

1. $o_l = \sigma_l (w_l \cdot i_l + b_l)$
2. i_0 is the input (feature)
3. o_L is the final output

- iii. compute final layer error
 - 1. compute gradient of cost function
 - a. example: quadratic cost function
 - b. $\nabla C = o_{\text{new}} - o_{\text{old}}$
 - 2. compute output error
 - a. $\delta_L = \nabla C \odot \frac{d\sigma_L}{dx}(o_L)$
 - iv. recursively compute prior layer errors
 - 1. $\delta_{l-1} = [(w_l)^T \cdot \delta_l] \odot \frac{d\sigma_{l-1}}{dx}(o_{l-1})$
 - v. update weight
 - 1. $\Delta w_l = \frac{\eta}{m} \delta_l (o_{l-1})^T$
 - 2. η learning constant, m layer inputs
 - vi. update bias
 - 1. $\Delta b_l = \frac{\eta}{m} \delta_l$
 - 2. η learning constant, m layer inputs
 - vii. repeat until satisfied with training
- d. training issues
- i. convergence
 - 1. cost function minimization
 - 2. problem solves general test cases sufficiently
 - ii. overfitting
 - 1. too much training fits the training set
 - 2. does not make for a general purpose solution
- e. mini-batches
- i. compute multiple forward stages
 - ii. cost function must satisfy summation condition
 - iii. computational complexity
 - 1. eliminates most of the backpropagation steps
 - 2. eliminates most of the weight and bias updates
 - 3. problem still remains in same class, just faster

7. Implementation

- a. intro to ocaml
 - i. applicative / functional programming
 - 1. lists
 - 2. first class functions
 - 3. higher order functions
 - ii. imperative features
 - 1. refs
 - 2. arrays

b. intro to blas / lapack

- i. blas 1
 - 1. vector - scalar
 - 2. $O(n)$ operations for $O(n)$ data
 - 3. memory bandwidth limited
- ii. blas 2
 - 1. matrix - vector
 - 2. $O(n^2)$ operations for $O(n^2)$ data
 - 3. memory bandwidth limited
- iii. blas 3
 - 1. matrix - matrix
 - 2. $O(n^3)$ operations for $O(n^2)$ data
 - 3. most able to be optimized
- iv. compare with lists in heap

c. lacaml

- i. brief overview of organization
- ii. what we are using it for

d. source

- i. math.ml
 - 1. functional
 - a. lists
 - b. heap
 - 2. imperative
 - a. arrays
 - b. fixed memory block
 - 3. performance and elegance tradeoffs
 - 4. binop versus full blas / lapack routines
- ii. types.ml
- iii. activation.ml
- iv. cost.ml
- v. layer.ml
- vi. network.ml
 - 1. three data structures
 - a. layer: $[\ell_1; \ell_2; \ell_3; \dots; \ell_L]$
 - b. forward: $[i; o_1; o_2; \dots; o_{L-1}]$
 - c. backward: $[\delta_1; \delta_2; \delta_3; \dots; \delta_L]$
- vii. utils.ml
- viii. time.ml

e. examples (including data)

- i. iris.ml
- ii. letters.ml
- iii. mnist.ml

- f. observations
 - i. issues training letters example
 - ii. mini-batch optimization and speed enhancements
 - 1. can treat mini-batch as matrix input instead of vector
 - 2. blas 3 optimizations matter
 - 3. significantly increased performance, same computational complexity

8. the future

- a. technique
 - i. half-precision floating point
 - ii. convolutional, recurrent networks
- b. gpus
 - i. baseline
 - 1. raspberry pi 2 b (used in demo)
 - 2. 98 mflops (single precision, one core)
 - ii. modern gpu
 - 1. nvidia tesla p100
 - 2. 18.7 tflops half-precision
- c. example libraries
 - i. accelerator
 - 1. cuda
 - 2. opencl
 - ii. machine learning
 - 1. tensor flow
 - 2. caffe
 - 3. theano